

Binary_Classification_using_Logistic_Regression_with_SparkMLib

September 6, 2023

Task 2 - Binary Classification using Logistic Regression

GOAL: The goal of this task is to build a machine learning pipeline including a classification model that predicts the **Attrition** (Yes or No) from the features included in the dataset (income, work years, education level, marital status, job role, and so on), which we used in the Lab 3 and Lab 4.

```
[0]: # IMPORTING THE NECESSARY LIBRARIES

# To convert categorical variables to numeric
from pyspark.ml.feature import StringIndexer

# To combine the feature columns into one single column
from pyspark.ml.feature import VectorAssembler

# For logistic regression
from pyspark.ml.classification import LogisticRegression

#For building the Pipeline
from pyspark.ml import Pipeline

# For checking the accuracy
from pyspark.ml.evaluation import BinaryClassificationEvaluator, \
    ↪MulticlassClassificationEvaluator

# For Hyperparameter tuning
from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
```

2.2 Loading the dataset and displaying the schema

```
[0]: # Loading the dataset
df1 = spark.read.format("csv").option("header", "true").load("dbfs:/FileStore/
    ↪shared_uploads/clvrashmika@gmail.com/EmployeeAttrition.csv", inferSchema = \
    ↪"true")

[0]: # Printing the dataset's schema
df1.printSchema()
```

```

root
|-- Age: integer (nullable = true)
|-- Attrition: string (nullable = true)
|-- BusinessTravel: string (nullable = true)
|-- DailyRate: integer (nullable = true)
|-- Department: string (nullable = true)
|-- DistanceFromHome: integer (nullable = true)
|-- Education: integer (nullable = true)
|-- EducationField: string (nullable = true)
|-- EmployeeCount: integer (nullable = true)
|-- EmployeeNumber: integer (nullable = true)
|-- EnvironmentSatisfaction: integer (nullable = true)
|-- Gender: string (nullable = true)
|-- HourlyRate: integer (nullable = true)
|-- JobInvolvement: integer (nullable = true)
|-- JobLevel: integer (nullable = true)
|-- JobRole: string (nullable = true)
|-- JobSatisfaction: integer (nullable = true)
|-- MaritalStatus: string (nullable = true)
|-- MonthlyIncome: integer (nullable = true)
|-- MonthlyRate: integer (nullable = true)
|-- NumCompaniesWorked: integer (nullable = true)
|-- Over18: string (nullable = true)
|-- OverTime: string (nullable = true)
|-- PercentSalaryHike: integer (nullable = true)
|-- PerformanceRating: integer (nullable = true)
|-- RelationshipSatisfaction: integer (nullable = true)
|-- StandardHours: integer (nullable = true)
|-- StockOptionLevel: integer (nullable = true)
|-- TotalWorkingYears: integer (nullable = true)
|-- TrainingTimesLastYear: integer (nullable = true)
|-- WorkLifeBalance: integer (nullable = true)
|-- YearsAtCompany: integer (nullable = true)
|-- YearsInCurrentRole: integer (nullable = true)
|-- YearsSinceLastPromotion: integer (nullable = true)
|-- YearsWithCurrManager: integer (nullable = true)

```

2.3 Splitting the dataset into training and testing sets & Displaying distribution of HourlyRate and Education

```

[0]: # Splitting the dataset into train and test dataframes
trainDF, testDF = df1.randomSplit([0.8, 0.2], seed=65)
print(trainDF.cache().count()) # Cache because accessing training data multiple
    ↪ times
print(testDF.count())

```

1204

```
[0]: # Checking the distribution of the 'HourlyRate' field in the training dataset
      ↪using the summary()
display(trainDF.select('HourlyRate').summary())
```

```
[0]: # Checking the distribution of the 'Education' field in the training dataset
      ↪using groupBY
display(trainDF.groupBy('Education').count().sort("count", ascending = False))
```

2.4 Feature Processing

```
[0]: # 2.4.1 - Selecting 5 categorical cols from the dataset
categorical_cols = ["Department", "EducationField", "Gender", "JobRole",
      ↪"MaritalStatus"]

# Covertng the above columns to numerical using stringIndexer
stringIndexer = StringIndexer(inputCols=categorical_cols, outputCols=[i +
      ↪"IndexedCol" for i in categorical_cols])

# 2.4.2 - Setting the Attritition Feature (Yes/No) as a label
# Converting to a numeric value
labelToNum = StringIndexer(inputCol="Attritition", outputCol="NewAttrititionCol")
labelToNum

#Applying this to the dataset
stringIndexerModel = stringIndexer.fit(trainDF)

labelIndexerModel = labelToNum.fit(trainDF)
```

```
[0]: # 2.4.3 and 2.4.4
# Combining the feature columns into a new single feature
numerical_columns = ["Age", "DailyRate", "Education", "DistanceFromHome",
      ↪"HourlyRate", "JobInvolvement", "JobLevel", "JobSatisfaction",
      ↪"MonthlyIncome", "YearsAtCompany", "YearsInCurrentRole",
      ↪"YearsWithCurrManager", "NumCompaniesWorked", "PerformanceRating",
      ↪"EnvironmentSatisfaction" ]

vector_assembler = VectorAssembler(inputCols=numerical_columns,
      ↪outputCol="features")
```

2.5 Defining the Model

```
[0]: # Defining the model for Logistic Regression
log_regression = LogisticRegression(featuresCol="features",
      ↪labelCol="NewAttrititionCol", regParam=1.0)
```

2.6 - Building the Pipeline

```
[0]: # Defining the pipeline based on the above created stages
pipeline = Pipeline(stages=[stringIndexer, labelToNum, vector_assembler,
    ↪log_regression])

# Defining the pipeline model
pipelineModel = pipeline.fit(trainDF)

# Apply the pipeline model to the test database
predDF = pipelineModel.transform(testDF)
```

2.6 (Cont.) - Displaying the Predictions

```
[0]: display(predDF.select("features", "NewAttritionCol", "prediction",
    ↪"probability"))
```

2.7 - Evaluating the Model

```
[0]: # Plotting the ROC curve
display(pipelineModel.stages[-1], predDF.drop("prediction", "rawPrediction",
    ↪"probability"), "ROC")
```

```
[0]: # Printing the area under the curve and the accuracy
binary_class_eval = BinaryClassificationEvaluator(metricName="areaUnderROC",
    ↪labelCol="NewAttritionCol")

print("Area under ROC curve: ", binary_class_eval.evaluate(predDF))

multi_class_eval = MulticlassClassificationEvaluator(metricName="accuracy",
    ↪labelCol="NewAttritionCol")

print("Accuracy: ", multi_class_eval.evaluate(predDF))
```

Area under ROC curve: 0.7369510015987963

Accuracy: 0.8157894736842105

2.8 - Hyperparameter Tuning

```
[0]: # Using ParamGridBuilder
parameterGrid = (ParamGridBuilder()
    .addGrid(log_regression.regParam, [0.01, 0.5, 2.0])
    .addGrid(log_regression.elasticNetParam, [0.0, 0.5, 1.0])
    .build())
```

```
[0]: # Using CrossValidator

#Creating a 3-fold CrossValidator
```

```
cross_validator = CrossValidator(estimator=pipeline,
    ↪estimatorParamMaps=parameterGrid, evaluator=binary_class_eval, numFolds=3)

# Running the cross validations to find the best model
cross_validator_model = cross_validator.fit(trainDF)
```

2.9 - Make predictions and Evaluate the model performance

```
[0]: cvPredDF = cross_validator_model.transform(testDF)

#Evaluating the Model performance
print("Area under ROC curve: ", binary_class_eval.evaluate(cvPredDF))
print("Accuracy: ", multi_class_eval.evaluate(cvPredDF))
```

Area under ROC curve: 0.7123107307439106

Accuracy: 0.8157894736842105

2.10 Use SQL Commands

```
[0]: # 2.10.1 Creating a temporary view of the predictions dataset
cvPredDF.createOrReplaceTempView("finalPredictions")
```

2.10.2 Displaying the predictions grouped by JobRole - Bar Chart

```
[0]: %sql
SELECT JobRole, prediction, count(1||2) as Count
FROM finalPredictions
Group By JobRole, prediction
Order By JobRole
```

Output can only be rendered in Databricks

2.10.3 Displaying the predictions grouped by Age - Bar Chart

```
[0]: %sql
SELECT Age, prediction, count(1||2) as Count
FROM finalPredictions
Group By Age, prediction
Order By Age
```

Output can only be rendered in Databricks

References:

1. Dr. Liao's Code Examples & Tutorials: Blackboard/Liao_PySpark_basic_databricks.html
2. PySpark: <https://spark.apache.org/docs/2.4.0/api/python/pyspark.html>